



INGQIQO EXECUTABLES (PTY) LTD

Company Profile

We build decision instruments — repeatable, documented tools that measure what an organisation's decisions cost, and whether they can be defended.

SEPTEMBER 2026
GAUTENG, SOUTH AFRICA

Meaning, made coherent.



AT A GLANCE

The short version

WHAT WE ARE

A South African decision intelligence practice. We do not sell software licences or advisory hours. We deliver instruments: a defined input set, a documented method, a checkable output, and a record of the decision that can be re-run and compared.

WHO WE SERVE

Mid-sized employers, municipal entities and public bodies, independent recruiters and small HR teams, boards and remuneration committees. The instruments are priced so a fifty-person business can use them, not only a listed bank.

WHERE WE STAND TODAY

Founder-operated and pre-revenue, delivering on pilot terms. Stated here rather than obscured. What can be inspected today — the method, the specimen outputs, the rate card — is listed on page 26.

WHY THE WORK HOLDS

Because the method, not the headcount, is the product. Every instrument is a scripted run of one documented cycle. A reviewer can check a result against the method rather than against an opinion — including an auditor who was never in the room.

WHAT A CLIENT RECEIVES

A figure with a pedigree

Observed, declared, published or estimated — with a range where estimated.

A stated limit

What the instrument deliberately does not measure or decide.

A named decision

The decision and its owner. We name it; we do not take it.

A review date

The date the recommendation can be judged against what happened.



I

Who we are

The position we hold, the method we work to, and the commitments a client can hold us to.



The position we hold

Rigorous decision analysis has long been delivered as bespoke advisory work — priced for organisations with large budgets and standing advisory relationships. The work was often excellent. It was out of reach for almost everyone else.

Those organisations carry the same costs. Undocumented rules. Approval queues nobody owns. Offers people cannot live on. Pay decisions nobody has checked. Decisions taken with no record of why. The costs are invisible, so nobody manages them.

Ingqiqo Executables exists to make those costs visible, at a price the organisation carrying them can actually pay. Not by cutting the rigour — by building the rigour into an instrument that can be run again, by someone else, and checked.

NOT THIS

A software licence

We sell no third-party platform and take no referral commission. A readiness assessment that concludes “no new system needed” is a valid outcome, and it costs us nothing to say so.

NOR THIS

Advisory hours

An engagement that leaves nothing behind but a document and a relationship has to be bought again next year at full price. Ours leaves a baseline the client owns.

THIS

An instrument

Defined inputs, a documented method, a checkable output, a stated limit. The second run costs a fraction of the first, and comparing the two is where the value actually lands.

Results, definitions and versions stay with the client, whether or not there is a next engagement. A baseline you cannot take with you was never a baseline.



The method — the Ingqiqo Disciplined Decision Cycle

Every instrument we deliver is a scripted run of one cycle, abbreviated IDDC. Six steps, in order, every time. The discipline is that no step may be skipped because a client is in a hurry.

1 Name the decision, and its owner

Before anything is measured. A diagnostic with no decision attached is a report, and reports do not change anything.

2 Define the unit

What exactly is being counted — one offer, one approval, one hour, one claim. A total built without this step is a number nobody can defend, and it will still print.

3 Measure what is there

The client's own tables and their own people, sampled anonymously. We do not substitute a sector average for a measurement we could have taken.

4 Calculate, and mark every figure

One results file. Each figure carries how it came to be — observed, declared, published or estimated — and a range where estimated. An estimate that looks like a measurement is the failure this step exists to prevent.

5 Render for each reader

One results file, several views — finance, operations, technology, employee relations, organised labour, staff, the executive. Different weightings, never different numbers.

6 Set the review date

Every significant recommendation carries the date it will be judged against what actually happened. This is the step that makes the second run worth buying.

WHY THE DISCIPLINE IS WRITTEN DOWN

A method held in one person's head cannot be audited, delegated or improved. IDDC is documented end to end and version-stamped on every output, which is what allows an accredited associate to run an instrument, a client to check it, and a successor supplier to pick it up. The full formal statement — its units, its memory horizons, its treatment of estimates — is set out in The Coherence Architecture, available under NDA where a technical evaluation requires it.



Calculation first — the model layer

A general-purpose language model produces the most plausible next word. That is a reasonable way to draft a paragraph and an unsound way to price a hire, size a pay gap or estimate a year of rework. Plausibility is not calculation. A board cannot audit a sentence that merely sounded right.

Jev is the posture that follows, and it governs our model layer: a figure is calculated on the client's own table, or it is not published. Jev decides which instrument runs, holds the balance between the two, and declines to answer where neither can carry the weight.

Language models keep a defined role. They read policies, render a stakeholder view and put findings into plain words. They never originate a figure. Every number in an Ingqiqo output was calculated, and traces to the row it came from.

THE POSTURE AND THE TWO INSTRUMENTS

Jev THE POSTURE

Holds the balance. Reads the shape of the client's table, selects the instrument, sets the range the estimate must carry, and refuses the question where the evidence cannot bear it.

TabDPT-Turbo AT VOLUME

Repeated tabular predictions that must run identically every time: offer decisions, churn and cost-of-acceptance estimates, benchmark variance.

TabICLv2 IN CONTEXT

First engagements and small tables: a twelve-month offer history, one department's time sample. Learns from the rows on hand, so a fifty-person client is not told to come back with more data.

All three are named and versioned on every output, alongside the method and parameter versions, so any result can be re-run and checked.

WHAT THIS MEANS FOR A PROCURING ENTITY

The engines are named and substitutable. The method, not the model, is the asset — so a replacement engine is a version increment and a re-run, not a rebuild, and no client is locked to a vendor we happen to use. Nothing in an Ingqiqo output depends on a model whose output we cannot show you the working for.



1.4

One substrate — the Decision Platform

Every instrument in this profile runs on the same platform, against the same definitions, under the same method. That is not an efficiency; it is the reason results are comparable — between two instruments, between two departments, and between this year and next.

FROM DAY ONE

IDDC was written before the first client

The method is not a framework we assembled from delivered engagements and named afterwards. It was documented first, and every instrument since has been built as a run of it. Nothing in our portfolio predates the discipline, which is why there is no legacy work that does not carry a pedigree.

ONE SUBSTRATE

The same platform under every instrument

A decision is expressed once — intent, evidence, assumptions, objectives, constraints, authority — and the engines act on that single record: tabular models supply predictions, a rules engine applies your policy, a solver tests feasibility, a simulator explores consequences. Each is evidence entering the decision; none of them is the decision. So a figure from one instrument sits beside a figure from another with nothing to reconcile.

CONSEQUENCE FOR THE CLIENT

Going deeper never means starting over

Because the substrate does not change between instruments, the baseline you own after the first engagement is the baseline the next one builds on. No re-collection, no new definitions, no second foundation to buy.

CONSEQUENCE FOR US

The work is consistent because the ground is

A practice that builds each engagement on whatever tooling suited it at the time cannot claim its results are comparable. Ours are comparable by construction. The platform also treats refusal as a valid answer — insufficient evidence, constraint unsatisfiable, out of distribution — because a system built to always answer will sometimes answer confidently and wrongly.

WHERE THE SUBSTRATE RUNS — THE LONG POSITION

NOW

Hosted, in-jurisdiction

Client data processed within the agreed jurisdiction, on infrastructure we rent and

NEXT

Owned control tier

Our own bare-metal cluster in colocation — the neutral ground every measurement

THEN

Private, privately-held models

Models we hold rather than call, so a client's figures

08



1.5

Eight commitments

These govern every client-facing output. They are testable by the client, which is what distinguishes them from a values statement.

01 One result, many views

No stakeholder receives a number the others cannot reproduce. Weightings differ between views; the underlying figures never do.

02 Every figure shows its source

Observed, declared, published or estimated — labelled on the page, with a range wherever estimated.

03 Every output shows its version

Method, engine and parameter versions printed on the document, so a result can be re-run and checked against the run that produced it.

04 Limits are stated

Each output says what it deliberately does not measure or decide. An instrument that claims to answer everything answers nothing reliably.

05 People decide

We name the decision and its owner. We do not take it, and no instrument of ours issues an automated rejection of a person.

06 Systems before people

Organisational diagnostics measure systems. They never measure, rank or name individuals, and are contractually prohibited from disciplinary or appraisal use.

07 Personal information is minimised

Only what a decision needs, under POPIA. Time is sampled anonymously and reported only for groups of five or more. Specimens are stripped.

08 Outcomes are checked

Every significant recommendation carries a review date. We would rather be measured and wrong than unmeasurable.



II

What we do

Six instrument families. Each stands on its own — you can stop after any of them and keep everything you have received.



The Employer Report

An employment offer is one of the most consequential decisions an organisation makes and one of the least documented. The Employer Report turns it from educated guesswork into a decision with a record — produced by the EETE engine on the CEDIE platform.

Cost of living the role requires

Housing, commute, debt service, schooling, professional clothing, technology — each line calculated and totalled, so an offer can be checked against what it actually takes to hold the job.

Cost of acceptance

Recruitment fee, onboarding, training, first-year total cost to company, time to full productivity — the full economic weight of the hire.

Role fit

Overlap against explicitly supplied skill lists. No black-box matching and no social-profile scraping — structured fields, reviewed by the recruiter.

Churn risk, with its drivers

A twelve-month estimate from seniority, economic margin, skills fit and experience band — each driver quantified and weighted, and each carrying its range.

Market position

Percentile positioning against the market data entered, with the source of that data declared on the page.

Process audit trail

Every stage from application to offer, with status, date and whether the record was explicit or inferred. A stage someone recorded and a stage the system deduced are not the same fact.

DECISION REASONS

The reasoning, set out for traceability. Thresholds crossed. Concessions noted. A verdict — justified, justified with conditions, or caution — that a remuneration committee can interrogate line by line.

SIGNAL CONFIDENCE

Robustness, the most sensitive input, and the result under an optimism stress test. When the same role on a second assumption profile gives a different answer, that is the argument for running it before the offer.

Stated limit. The report is a calculation on supplied inputs, not a prediction of what any individual will accept. It is not financial, tax, legal or employment advice, and it does not decide the hire.



The Organisational Debt Diagnostic

Decisions taken years ago are still being paid for — in time spent reconstructing what was agreed, in claims reopened because two teams define a metric differently, in people leaving above what the sector does. The diagnostic prices that. It is delivered on the Organisational Debt Ledger, which is what makes the second run a comparison rather than a fresh opinion.

Debt-attributable drain

Gross time spent compensating for past decisions, less a normal coordination floor — typically eight per cent of loaded payroll. Reported as a rand figure and as a share of payroll, so it can sit in a budget conversation.

Where the debt goes

Archaeology — finding, asking, reconstructing. Rework — claims reopened after a definition dispute. Waiting — approvals, information, decisions. Excess meetings — twenty-two people, ninety minutes, weekly, reading a dashboard aloud. Each quantified annually.

Excess attrition and risk

Turnover above the sector reference by role, with replacement cost. Risk exposure stated probability-weighted rather than worst-case, because a worst-case number gets discounted and ignored.

The debt register

Thirteen or more named debts, each with a principal, an interest rate, a structural rank, its downstream dependencies, and a disposition: repay, tolerate, defer, or convert to a managed option. Tolerating a debt deliberately is a legitimate answer — an untracked one is not.

WHERE THE FINDING USUALLY IS

The financial ranking and the structural ranking normally disagree about which debt matters most, and the disagreement is the finding. Net present value says repay the loudest item. The dependency graph often says that item is a symptom of something upstream — and that repaying only the leaves means paying the same interest again next year.



Decision governance and maturity

Most organisations cannot say who is entitled to decide what, on what evidence, accountable to whom. That is not a document problem. It is why the same decision gets taken twice, differently, in two parts of the business.

INSTRUMENT

Decision maturity assessment

A structured diagnostic across the dimensions that determine whether a decision system works: who holds authority, what evidence is required, how exceptions are handled, what is recorded, and whether anyone checks the outcome. Scored, so the second run is comparable.

Produces a scored profile, a gap list ranked by consequence, and a named owner per gap.

INSTRUMENT

Governance review and Decision Canvas

A working session that puts a recurring decision on one page: the decision, its owner, the evidence required, the rule in force, the exception path, and the review date. The Canvas is the artefact the organisation keeps and maintains itself.

Produces a Canvas per decision class, and a gap assessment against our decision governance standard.

INSTRUMENT

Analytics readiness assessment

Whether the organisation's data is in a state where decision instruments can be built on it at all — definitions, coverage, ownership, and whether a derived number is distinguishable from an observed one.

"No new system needed" is a valid and common outcome. We sell no platform, so it costs us nothing to reach it.

FOR BOARDS AND PUBLIC BODIES

The audit answer

When an audit asks how a decision was reached, the usual answer is a reconstruction. These instruments produce the record at the time of the decision instead — the decision, the evidence in force, the rule applied, the owner, the date. Our governance work supports King IV obligations; it does not replace them, and every output says so.

A governance review that produces a policy nobody maintains has failed. Every Canvas has an owner and a review date, and the organisation is trained to keep it without us.



2.7

Sector applications

The method does not change by sector. What changes is which decision is worth instrumenting first.

SECTOR	THE DECISION WORTH INSTRUMENTING FIRST	STARTING INSTRUMENT
Municipal and public entities	Appointments and procurement recommendations that must survive an audit two years later, where the record was never made at the time	Governance review and Canvas
Mining and heavy industry	Scarce-skill appointments at remote operations, where an offer that does not cover the real cost of living there fails in month four	Offer audit, then Employer Report
Financial services	Remuneration decisions under committee scrutiny, and metric definitions that differ between divisions	Pay equity review
Manufacturing and logistics	Approval queues and undocumented exception handling that nobody owns but everybody works around	Organisational Debt Diagnostic
Agriculture and co-operatives	Seasonal labour and input timing decisions taken annually with no record of what last season's decision actually cost	Governance review
Professional services and recruiters	Placement advice given on instinct, and a client asking why the last one did not hold	Employer Report subscription
Education and health	Staffing and rostering decisions where the human load is obvious to everyone and measured by nobody	Organisational Debt Diagnostic

WHAT WE WILL NOT CLAIM

We do not hold delivered engagements in every sector above. The table states where each instrument applies and why, not where it has already been run. Our pilot record is on page 27, stated plainly, and a sector where we have no experience will be named as such in any proposal.



III

How we work

The engagement itself, the capacity behind it, and what happens to your people's data while we hold it.



3.1

An engagement, end to end

A five-week diagnostic, as it actually runs. Shorter instruments compress this; none of them skip a stage.

WEEK 1

Scope

The decision is named and its owner identified. Scope, price and exclusions are fixed in writing. Staff and, where present, organised labour are briefed on what is being measured and what it will not be used for — before any data is touched.

WEEKS 2-3

Measure

Anonymous time sampling, system extracts and structured interviews. Sampling is voluntary and reported only for groups of five or more. Where a number cannot be measured we say so and estimate it with a range, rather than substituting a sector average silently.

WEEK 4

Compute and check

One results file. Figures calculated on your tables, each marked with its pedigree and range, method and engine versions stamped. Internal review is against the documented method, not against a colleague's judgement.

WEEK 5

Render and hand over

Stakeholder views rendered from the one file — finance, operations, technology, employee relations, organised labour, staff, and an executive decision brief. Different weightings, identical numbers. The decision brief states what must be decided, by whom, and by when.

AFTER

Review

Data is returned or destroyed on the agreed schedule. The baseline, the definitions and the versions stay with you. On the review date the instrument can be re-run at a fraction of the first cost, and the comparison is the point — it is the only evidence that anything improved.

WHAT WE NEED FROM YOU

A named decision owner with authority. Access to the relevant extracts. Roughly six to ten hours of your people's time across the engagement. Permission to brief staff honestly about what is being measured.

WHAT YOU WILL NOT BE ASKED FOR

A data warehouse, a platform purchase, or a full population census. If your data is not ready, a readiness assessment says so — and "not yet worth building" is an outcome we are paid to deliver.



People, data and protection

A diagnostic that measures how an organisation works is measuring people's working lives. That deserves a stated position, not a privacy clause in the annexures.

SYSTEMS, NEVER INDIVIDUALS

Organisational diagnostics measure processes and decision rules. They name no individual and rank no employee. Time is sampled anonymously and reported only for groups of five or more, so no reading can be traced back to a person.

PROHIBITED USES, CONTRACTUALLY

Diagnostic data may not be used in disciplinary, incapacity, appraisal or retrenchment-selection decisions. This is a clause in the engagement terms, not a statement of intent, and it survives the end of the engagement.

MINIMISATION UNDER POPIA

We collect only what the named decision requires. Specimens and samples are stripped of identifying detail. Data is returned or destroyed on the schedule set in the engagement terms, and the sole director is the registered information officer.

WHERE LABOUR IS ORGANISED

We recommend a joint governance body with audit rights over how the data is used, and staff receive their own plain-language view of the findings. A diagnostic that only management ever reads generates resistance the next time, and deservedly so.

THE STAFF VIEW

Every organisational diagnostic produces a view written for the people whose work was measured. It states what was measured, what was found, what will change, and what the data will never be used for. It is rendered from the same results file as the executive brief, with the same figures – which is the whole point of the commitment on page 08.

An instrument that measures a workforce and reports only upwards has told half the truth to half the people. We do not deliver those.



Compliance and governance

OBLIGATION

OUR POSITION

POPIA

Personal information minimised by design; group-of-five reporting floor; stripped specimens; a stated return-or-destroy schedule; registered information officer. PAIA manual and processing register maintained.

Employment Equity Act

Where an instrument assesses an individual, any psychometric input must meet the Act's requirement that tests be valid, reliable, fairly applied and unbiased. No instrument of ours triggers an automated rejection of a candidate.

Labour Relations Act

Diagnostic data is contractually prohibited from disciplinary, incapacity, appraisal and retrenchment-selection use.

King IV

Our governance instruments support King IV obligations. They do not discharge them, and every relevant output states this limit on its face.

Conflict of interest

We resell no third-party software and take no referral commission. Declared in every proposal. A readiness assessment concluding "no platform needed" carries no revenue penalty for us — which is what makes the finding worth something.

Data residency

Client data is processed within the agreed jurisdiction. Residency is an execution guarantee — that data does not leave. It is stated separately from the governance guarantee, which is that the number means what it says. The two are different promises and we do not conflate them.

Professional standing

We are not lawyers, auditors, actuaries or tax practitioners and make no such claim. Where an engagement requires one of those opinions we name the registered professional who will provide it.

AUDITABLE WITHOUT OUR COOPERATION

Because every output prints its method version, parameter set and the pedigree of each figure, a deliverable can be checked after the engagement closes by someone who was never in the room — an internal auditor, a commissioner, or a successor supplier. That is a deliberate design choice, and it is the strongest assurance a small supplier can offer.



IV

Working with us

Where to start, how partners work with us, who does the work, and every document an evaluation committee will ask for.



Partners and accredited associates

Because the method is documented rather than personal, other firms can deliver it. That is the intent: a standard that only its author can apply is not a standard.

ROUTE 1

Accredited associate

An independent practitioner who delivers facilitation, interviews and workshops on our engagements, to the documented method and under our review. Engaged per engagement, not retained. Certification at practitioner level is the qualification; no associate signs a finding – the company does.

ROUTE 2

Accredited implementation partner

A firm accredited to deliver named instruments to its own clients under licence, on a revenue-share basis. Requires certified staff, an annual quality review of delivered work, and acceptance of the eight commitments on page 08 as binding terms.

ROUTE 3

Referral and joint bid

For tenders where our instruments are one component of a larger scope, or where a prime contractor needs decision analysis it does not deliver itself. We are content to be a named subcontractor with our scope, price and ceiling stated separately in the bid.

ROUTE 4

Academic and standards collaboration

Co-developed curriculum with universities, and joint submissions to standards and professional bodies. We publish the method rather than protect it, so collaboration on the standard itself is welcome and expected.

WHAT A PARTNER GETS

The documented method, the templates, training and certification, the benchmarking reference set, and a named instrument they can quote against a published rate card.

WHAT A PARTNER ACCEPTS

The eight commitments, the prohibited-use clauses, an annual review of delivered work, and the rule that an instrument delivered outside the method may not carry our name.



INGQIQO EXECUTABLES (PTY) LTD

Meaning, made coherent.

A decision you cannot explain is a decision you cannot defend,
improve, or repeat.

Every instrument in this profile exists to do one thing: leave behind a record of a decision that someone else can check. The figure with its pedigree. The limit, stated. The owner, named. The review date, set. That is the whole product, and it is why the second run matters more than the first.

We are a small practice and we say so. What we offer instead of scale is that our work can be audited without our cooperation — by your internal auditor, by a commissioner, or by whoever comes after us. A supplier who builds that in has already answered the only question that matters.