



INGQIQO EXECUTABLES (PTY) LTD

Meaning, made coherent.

ISSUED AT THE START OF EVERY ENGAGEMENT

Engagement Pack

How we will work with you, what we will and will not do, who decides what, and the answers to the questions you would otherwise have to ask us one at a time.

Engagement: [client] · [scope]

Issued: [date] · Pack version 1.0

NOT A LEGAL INSTRUMENT



READ THIS FIRST

What this pack is, and what it is not

IT IS

Our conduct, written down in advance

A statement of how we work, what we commit to, what we refuse, who decides what, and how disagreement is handled. It is written to be quoted back at us. If our conduct departs from it, this pack is the evidence.

IT IS NOT

A contract, and not legal advice

It creates no obligation on either party and overrides nothing. Where it and a signed agreement differ, the agreement governs and we will say so rather than argue. Nothing here is legal, tax, audit or actuarial advice, and we are not registered to give any of it.

SO WHY ISSUE IT AT ALL

Because most of what goes wrong in an engagement is not a breach of contract. It is a mismatch of expectation that nobody wrote down: what "done" means, who was entitled to decide, whether a number was measured or estimated, what the data would be used for, what happens when scope grows. A contract is drafted to allocate blame once that has already happened. This pack exists to prevent it – and it is deliberately the kind of document a procurement officer, a company secretary and a shop steward can each read in full.

HOW THE PACK IS USED

AT FIRST CONTACT

Issued before a proposal, so you can decline us on the terms rather than on the price.

AT KICK-OFF

Walked through with the decision owner, and with staff or labour representatives where people are measured.

DURING DELIVERY

The reference for the questions on pages 12 to 16. If the answer here is wrong, raise it.

AT CLOSE

The checklist against which the engagement is reviewed, by you.

Every commitment in this pack is one we would be willing to have published to the people it affects. That is the single test we apply to our own conduct, and it is the reason this document can be handed to anyone in your organisation without editing.



Contents

A · THE ENGAGEMENT

What we are doing, and what done means	04
Who decides what	05
The week-by-week shape	06
What we need from you	07

B · THE COMMITMENTS

What we warrant, and what we do not	08
Your people, your data	09
Theory of welfare, for signature	10
Money, time and how scope changes	11

C · THE ANSWERS

Questions your executive will ask	12
Questions your staff will ask	13
Questions procurement and audit will ask	14
The platform beneath the work	15
When we disagree	16

D · CLOSE

Handover, and what you keep	17
The engagement review, run by you	18

COMPANION DOCUMENTS

Seeing Clearly — the instruments and their limits, in plain language. The Currency of Craft — how the work is priced and why. Company Profile — capability and compliance position. The Coherence Architecture — the method in full, under NDA where a technical evaluation needs it.



A.1

What we are doing, and what done means

Completed for this engagement at kick-off. Nothing below is boilerplate — if a line is blank when work starts, work has not started.

The decision this serves	[named, in one sentence]
Decision owner — named, with authority	[one person, not a committee]
The kind of work it is	[fixing · migrating · building · automating · evaluating · productionalising]
The unit being counted	[one offer · one approval · one hour · one claim]
Deliverables, itemised	[listed, each acceptable or not on its own]
The acceptance test	[what you will check, and against what]
Explicitly out of scope	[named, so it cannot drift in]
Baseline agreed before we start	[stated at a named grain, with its pedigree]
Review date	[when the recommendation gets judged against what happened]

WHY THE TYPE IS NAMED

Most scope disputes are really disputes about which kind of work was commissioned. A fix is accepted against the fault; a productionalisation against behaviour under failure. Agreeing that now is the cheapest risk reduction available to either of us.

WHY THE BASELINE COMES FIRST

A baseline established after an intervention is a negotiating position, not a measurement. Agreed beforehand, it is the only thing that lets either of us say honestly, at the review date, whether anything moved.



A.2

Who decides what

We name the decision and its owner. We do not take the decision. This table makes that boundary explicit, because it is the one most often blurred by a supplier who wants to be indispensable.

MATTER	DECIDED BY	OUR ROLE
What the organisation does about a finding	You	Name the decision, cost the options, state the trade-off
Whose interest the design serves where interests conflict	You, at named-executive level	Surface the conflict rather than settle it in a template
Whether a figure is fit to publish internally	You	Mark its pedigree and range so you can judge
Whether the method was followed	Either party, against the standard	Publish the standard and the version stamps
Whether a figure is arithmetically correct	Us – and our error to fix at our cost	Re-derive on request
Whether the work can be delivered in the window	Us, before award, in writing	Decline or reschedule rather than absorb
Whether an intervention is withdrawn	Its named owner, on your side	Build it to be reversible, and say how
A legal, tax, audit or actuarial conclusion	A registered professional	Name who should give it; we do not

THE ONE THING WE ASK YOU TO HOLD

A single named decision owner with actual authority. Not a committee, not a liaison. An engagement with no owner produces a report instead of a decision, and a report is the most expensive way to change nothing.



A. 3

The week-by-week shape

A five-week diagnostic as it actually runs. Shorter instruments compress this; none of them skip a stage, and we will tell you which stage we are in without being asked.

WEEK 1		
Scope	Page 4 completed and signed off. Decision and owner named. Baseline fixed. Staff and, where present, labour representatives briefed on what is being measured and what it will never be used for — before any data is touched.	You receive: the completed page 4, and the welfare statement on page 10
WEEKS 2-3		
Measure	Anonymous time sampling, system extracts, structured interviews. Sampling is voluntary and reported only for groups of five or more. Where something cannot be measured we say so and estimate it with a range, rather than substituting a sector average quietly.	You receive: a weekly written update, and immediate notice of anything that changes scope
WEEK 4		
Compute and check	One results file. Figures calculated on your own tables, each marked with its pedigree and range, with method and engine versions stamped. Internal review is against the published standard, not against a colleague's opinion.	You receive: notice of any finding that will be uncomfortable, before it is written up
WEEK 5		
Render and hand over	Views rendered from the one file — finance, operations, technology, employee relations, labour, staff, and an executive decision brief. Different weightings, identical numbers. The brief states what must be decided, by whom, by when.	You receive: every view, including the one written for the workforce
AT 90 DAYS		
Review	The outcome review against the Week 1 baseline. Included unless you opt out. We report what moved, what did not, and where our own estimate was wrong — that last one being the part that makes the rest credible.	You receive: the review, and a continuation proposal you are free to ignore

Nothing renews automatically. Each phase ends at a decision point. A rollover we did not have to earn would contradict every commitment in Part B.



A. 4

What we need from you

Short list, deliberately. An engagement that requires heroic effort from your team to keep alive was scoped wrongly.

WE NEED

A decision owner with authority

One person, named on page 4, who can act on the finding.

Access to the relevant extracts

The tables the instrument actually needs, not a data warehouse.

Six to ten hours of your people's time

Across the whole engagement, in scheduled slots, not on demand.

Permission to brief staff honestly

What is measured, and what it will never be used for. If this is refused, we will not take the work.

An answer on whose welfare counts

Page 10 cannot be completed by us alone. Where interests conflict, that is your call to make and record.

WE WILL NOT ASK FOR

A platform purchase

We resell nothing and take no commission. If your data is not ready, that is a finding, not an upsell.

A full population census

Sampling is sufficient and is better for the people sampled.

Identified personal records

We ask for the minimum the named decision requires, and no more.

A reference before we have earned one

We will ask at the 90-day review, once there is something to attest to.

Exclusivity, or a standing retainer

The baseline is yours; you may commission the next run from anyone, including a competitor.



B.1

What we warrant, and what we do not

The right-hand column is the one that matters. A commitment without stated exclusions is a sales line that gets disclaimed in an annexure nobody read.

WE STAND BEHIND

The arithmetic

Every figure re-derivable from the inputs supplied at the stamped versions. Our error is ours to correct at our cost.

That the method was followed

Each output meets the published standard item by item, checkable by anyone.

That the pedigree is honest

Nothing estimated is shown as observed. Estimates carry ranges. Thin sources are labelled thin.

That it reproduces

Same inputs, same versions, same answer. If not, that is a defect and we treat it as one.

That the data is handled as agreed

Minimised, reported only in groups of five or more, returned or destroyed on schedule, never used for a prohibited purpose.

WE DO NOT STAND BEHIND

Any commercial outcome

A finding does not guarantee a saving. What happens next is a decision we name but do not take.

Inputs we did not measure

A figure you supply is marked declared. We warrant our use of it, not its accuracy.

Individual behaviour

An estimate is a probability over a population. No output of ours is evidence about an identified person.

Legal, tax, audit or actuarial conclusions

We are not registered to give them. Where one is needed we name who should.

Compliance itself

Our governance work supports obligations under King IV and the employment statutes. It does not discharge them.

Both columns appear in our engagement terms in these words. Nothing on the left is softened there; nothing on the right is added there. Read this page and you have read the commercial substance.



Your people, your data

This page is written to be handed to the people whose work is measured, unedited. If any of it would embarrass us to show them, it should not be here.

SYSTEMS, NEVER INDIVIDUALS

Organisational diagnostics measure processes and decision rules. They name no individual and rank no employee. Time is sampled anonymously and reported only for groups of five or more, so no reading traces back to a person.

PROHIBITED USES

Diagnostic data may not be used in disciplinary, incapacity, appraisal or retrenchment-selection decisions. This is a clause in the terms, not an intention, and it survives the engagement.

MINIMISATION

We collect only what the named decision requires, under POPIA. Specimens are stripped of identifying detail. Data is returned or destroyed on the stated schedule. Our sole director is the registered information officer.

WHERE LABOUR IS ORGANISED

We recommend a joint body with audit rights over how the data is used, and staff receive their own plain-language view of the findings. A diagnostic only management ever reads earns resistance next time, deservedly.

TWO DISTINCTIONS WE HOLD, AND WILL NOT BLUR

Observed is not inferred

A recorded fact and a value our models filled in are different facts, and they are marked differently on every page. A behavioural estimate never silently becomes a measurement — that is the single most common way this kind of work does harm.

Predicted is not intended

An estimate that a pattern is likely is not a statement that anyone chose it, and not a judgement about anyone's conduct. Where an outcome emerged from how the system is arranged rather than from a decision, we say so.



Theory of welfare, for signature

Completed at kick-off and countersigned. We redesign the context in which decisions are made — defaults, sequencing, evidence, thresholds, templates. That is an intervention in how people choose, so whose interest it serves is stated rather than assumed.

Primary interest	[stated — e.g. long-term value of the organisation]
Secondary interests	[stated — e.g. employees affected by the process; customers or citizens affected by its outcomes]
Where they conflict	We will say so and escalate the trade-off to [named executive] for decision. We will not settle it silently through design.
Freedom	Every default has a clear, low-cost way to choose otherwise. Friction that serves the institution against the chooser is out of scope.
Publicity	We will not implement anything we would be unwilling to disclose to those it affects. Each intervention is documented in plain language.
Endorsement	We steer only toward choices the affected decision-makers would endorse on reflection, and we check that with them rather than assume it.
Capability	Each redesign builds your people's own decision capability, so sustaining it does not depend on us.
Measurement	Each intervention has a baseline, a stated expected effect and a review date. We report results — including null and negative ones.
Contestability	Anyone affected may challenge an intervention through [named channel]. Challenges are logged and answered.
Reversibility	Every intervention has a named owner and can be withdrawn if the evidence or the affected stakeholders do not support it.
Register	Every intervention is entered in a register held by you: purpose, whose interest it serves, the opt-out path, expected and measured effect, owner, review date, status. An intervention whose assumptions have expired is withdrawn rather than left running.

Signed for Inggiqo Executables

Signed for the client

Date

ON THE SIZE OF THE CLAIM

Interventions of this kind, measured honestly in real organisations, move outcomes by low single-digit percentage points — far below what published studies report. We promise measurement, not magnitude. On a large enough flow a small honest percentage is a substantial number; a large dishonest one is a dispute.



Money, time and how scope changes

PRICE

Quoted from the published rate card

Never priced by what you appear able to pay. Two clients buying the same instrument in the same month sit in the same band, and either may ask where in it they sit and why.

INVOICING

Half on commencement, balance on the brief

The deposit covers the run, because the cost base is time rather than materials. Any client may request the task-level time record behind an invoice and receive it.

TIME RECORDS

Daily, at task level, producible

Each entry names what was done and what was produced. Our exploratory work, learning and admin are ours to carry and never appear on your invoice.

SCOPE CHANGE

Flagged the week it happens

When a task exceeds its estimate we tell you immediately, with the revised number and the reason. Absorbing it silently hides a change from the person entitled to decide about it.

WHAT THE ESTIMATE ALREADY INCLUDES

Profiling and investigation — the fault is rarely where it was reported	15–25%
Data quality remediation — discovered during execution, not before	10–20%
Testing and validation — the first thing cut under pressure, and the most dangerous	20–30%
Documentation — otherwise deferred and never completed	10–15%
Change and communication where users are affected	10–20%
Rework provision — assumed to be zero everywhere else; it is never zero	10–15%

Why we show you the allocation. An estimate that omits these is not cheaper — it is the same price arriving later as a variation order. You can now compare our quote against anyone else's on the same basis, and ask them which of the six they provided for.



Questions your executive will ask

Our answers, in advance, so your sponsor can answer them without waiting for us to be in the room.

“We already have dashboards, BI and AI. What is different?”

Dashboards present information. What we produce is a record of the transition from evidence to authorised action — intent, assumptions, evidence and its pedigree, alternatives, constraints, who decided, and what actually followed. The differentiator is not a better chart; it is that the decision itself becomes inspectable afterwards.

“Isn't this just change management?”

It overlaps and goes lower. Change management asks how people adopt a change. We ask how this organisation repeatedly produces choices and commitments under uncertainty — and then alter the templates, thresholds and sequencing that produce them.

“Our decisions are too unique to standardise.”

We do not standardise the judgement. We standardise the scaffolding around it: what evidence is required, how options are structured, who decides, what gets recorded, and how the outcome is checked. Discretion is preserved inside explicit constraints — that is the whole design.

“What is the return, in a number I can defend?”

We cost the friction in the flow we are given — delay, rework, error rate multiplied by stakes — against a baseline agreed before we start, and report what moved at ninety days. We will not quote you a percentage from someone else's study. On a flow of any size, a small measured improvement is a large number; an unmeasured large one is a dispute waiting to happen.

“Who loses discretion if this works?”

The right question, and we will raise it if you do not. Making decision rights explicit exposes who actually decides, on what evidence. Some resistance to this work is not about priorities — it is about that exposure. We will name where we expect it, and we will not pretend the engagement is politically weightless.



Questions your staff will ask

Answered plainly, and meant to be circulated as written. These are the questions people ask each other rather than us, which is exactly why they need answers on paper.

“Is this a headcount exercise?”

No, and we have removed our own incentive for it to become one: where a gain would come from reducing headcount, we decline any share of it and charge a fee instead. We measure systems, not people, and we hold no commercial interest in fewer of you.

“Will my manager see how I spend my time?”

No. Sampling is voluntary and anonymous, and nothing is reported for a group smaller than five. There is no individual output, no ranking and no per-person figure — not in the executive brief and not anywhere else.

“Can this be used against me later?”

Contractually, no. Diagnostic data may not be used in disciplinary, incapacity, appraisal or retrenchment-selection decisions, and that clause survives the end of our engagement. If you are told otherwise, this pack is the document to produce.

“What if the findings are wrong about my work?”

Then say so through the channel named on page 10. Challenges are logged and answered, not absorbed. Every figure shows whether it was observed, declared, published or estimated, so a disagreement can be settled by re-running the instrument rather than by argument.

“Will we hear what was found, or only management?”

You will hear it. Every organisational diagnostic produces a view written for the people whose work was measured — what was measured, what was found, what will change, and what the data will never be used for. It is rendered from the same file as the executive brief, with the same numbers.

“If something changes and it makes the job worse, then what?”

Every change we make has a named owner and can be withdrawn. That is a design requirement, not a courtesy: if the evidence or the people affected do not support an intervention, it comes out. Friction that serves the institution against the person doing the work is the thing we are hired to find, not to install.



Questions procurement and audit will ask

QUESTION

ANSWER

Is the supplier registered, tax compliant and verifiable?

Yes. Registration, VAT, tax compliance PIN, B-BBEE and supplier database details are in the Company Profile, with certificates attached to any submission.

Can the deliverable be audited independently?

Yes, and without our cooperation. Every output prints its method and engine versions and the pedigree of each figure, so an internal auditor or successor supplier can check it after we have gone.

What happens if the key person is unavailable?

The method is documented rather than held in one head, and accredited associates are trained against it. Professional indemnity and key-person cover are placed before any public-sector award.

Is there a conflict of interest?

Declared in every proposal: we resell no software and take no referral commission. A finding of "no new system needed" carries no revenue penalty for us.

Where is our data processed, and for how long?

Within the agreed jurisdiction, on the platform described on page 15, returned or destroyed on the schedule in the terms. Residency is stated separately from the governance guarantee; they are different promises.

Is the supplier POPIA compliant?

Minimisation by design, a group-of-five reporting floor, a registered information officer, a PAIA manual and a processing register.

Is any decision automated?

No. No instrument of ours issues an automated rejection of a person, and every recommendation has a human owner named on page 5. Where a model contributed a value, that value is marked as model-derived on the page.

Are there comparable past engagements?

We are pre-revenue and delivering on pilot terms, and we say so. Where mandatory criteria require a trading history we do not hold, we decline rather than construct an answer.

Is the pricing transparent?

One published rate card, quoted from nowhere else, with the six estimate allocations on page 11 disclosed so quotes can be compared on the same basis.

One thing we will not do. We will not answer a mandatory criterion we cannot meet by reinterpreting it. If that costs us the award, it costs us the award — a supplier who overstates its position once cannot be relied on for a finding.



The platform beneath the work

You are not buying software and we are not asking you to adopt any. But every instrument runs on one substrate — the Decision Platform — and its design is why the commitments in Part B can be made at all. This page is what a technical evaluator needs.

ONE RECORD, MANY ENGINES

Each decision is expressed once — intent, evidence, assumptions, objectives, constraints, authority — and that single record is what the engines act on. Tabular models supply predictions, a rules engine applies your policy, a solver tests feasibility, a simulator explores consequences. None of them is the decision; each is evidence entering it.

SUBSTITUTABLE BY DESIGN

Engines are named and versioned on every output, and any of them can be replaced without touching the record. That means no lock-in to a vendor we happen to use, and it is why a replacement is a version increment and a re-run rather than a rebuild.

CALCULATION FIRST

Figures are computed on your own tables — purpose-built tabular models, not general-purpose language models. Language models draft prose from the results file and never originate a number. A figure that was calculated can be re-derived, which is the precondition for warranting it.

REFUSAL IS A VALID OUTPUT

The platform can return insufficient evidence, constraint unsatisfiable, or out of distribution — and does. A system that always produces an answer is a system that will sometimes produce a confident wrong one, and we would rather hand you a refusal than a plausible number.

FOUR DISTINCTIONS THE PLATFORM ENFORCES, NOT MERELY OBSERVES

Observed ≠ inferred

A model-filled value is marked as one, per cell, and never inherits the standing of a measurement.

Predicted ≠ intended

A likely pattern is not evidence that anyone chose it.

Optimal ≠ desirable

A solver result is an option, subject to the welfare statement on page 10.

Emergent ≠ decided

Outcomes produced by how the system is arranged are reported as such, not attributed to a person.

Where it runs. Today, on infrastructure we rent within the agreed jurisdiction and can attest to. Our stated direction is an owned control tier and privately-held models, funded from surplus — set out in the Business Plan. That is a capital direction, not something priced into this engagement. The full method sits in The Coherence Architecture, available under NDA where a technical evaluation requires it.



When we disagree

Written before there is a disagreement, which is the only time it can be written honestly.

IF YOU DISPUTE A FIGURE

We re-derive it in front of you at the stamped versions. If our arithmetic is wrong it is our error and our cost. If the figure is right but rests on an input you supplied, it is marked declared and we will show you where that shows on the page.

IF YOU DISPUTE A FINDING

We do not soften a finding to preserve a relationship. We will state the disagreement in the deliverable, in your words and in ours, and leave the reader to judge. A finding we withdrew under pressure would make every other finding worthless.

IF WE ARE LATE

You hear it the week we know, not at the deadline, with a revised date and the reason. We would rather deliver late having told you early than on time having cut the testing.

IF YOU WANT TO STOP

Stop. You pay for work done to that point, you keep everything produced, and you owe no explanation. There is no termination penalty and nothing renews on its own.

IF WE SHOULD STOP

We will say so. If the data cannot support the question, if the decision owner has no authority to act, or if we are being asked to produce a number to justify a decision already taken, we will end the engagement and invoice only what was done. That is a commitment against our own revenue, which is the only kind worth printing.

Escalation. First to the decision owner named on page 4. Then to our principal, who is accountable for every finding issued in the company's name and does not delegate that. There is no third tier, because there is no one behind us to hide behind.



D.1

Handover, and what you keep

The test of this engagement is whether you could run it again without us. Everything below is yours, by default, without asking.

YOU KEEP

The baseline and its definitions

At the grain agreed in week one, with every pedigree mark intact.

The results file, not just the report

The thing every view was rendered from, so you can render your own.

Method and engine versions

Enough to re-run and compare, by us or by anyone else.

The intervention register

Every change, its purpose, whose interest it serves, its opt-out path, its owner and its review date.

A trained internal owner

Someone on your side who can maintain what we built. Where a redesign is involved this is a deliverable, not a courtesy.

WE RETAIN

The method itself

Ours, and licensed to you for your own use of the instrument — not sold with the engagement.

Anonymised benchmark contributions

Only with your written agreement, stripped, and never in a form that identifies you. You may decline and it changes nothing about the engagement.

Nothing else

Your data is returned or destroyed on the stated schedule. We keep no working copy, and we will confirm the destruction in writing.

WHY WE HAND OVER THIS MUCH

Because a baseline you cannot take with you was never a baseline, and because the second run is worth more to you than owning you is worth to us. You may commission the comparison from a competitor. We would rather compete for it than hold it by default.



D.2

The engagement review, run by you

Score us against this pack at the ninety-day review. Twelve questions, each answerable yes or no by someone who was there. We will ask for it whether or not you offer it.

- | | | | |
|----|--|----|--|
| 01 | Was page 4 completed and signed before work began? | 02 | Was the baseline fixed before any intervention? |
| 03 | Were staff briefed honestly, before data was touched? | 04 | Did every figure show whether it was observed, declared, published or estimated? |
| 05 | Did every output print its method and engine versions? | 06 | Did each deliverable state what it deliberately did not measure? |
| 07 | Did every stakeholder view carry the same numbers? | 08 | Did the workforce receive its own view of the findings? |
| 09 | Was scope growth flagged the week it happened, in writing? | 10 | Was the welfare statement completed, signed, and honoured? |
| 11 | Did we report where our own estimate was wrong? | 12 | Could you re-run this next year without us? |

IF THE ANSWER TO ANY OF THESE IS NO

Then we failed a commitment we published in advance, and this page is the evidence. Tell us, in writing, and we will record it in the review alongside your words. A supplier who publishes a checklist and then argues with the score has published nothing.

This pack is not a contract, and it binds us more tightly than one — because a contract is argued in private and this was published to everyone it affects.

Contact details — to complete

PACK VERSION 1.0
GAUTENG, SOUTH AFRICA